

Three training sources → *model* | *held-out source* → *calibrate + test*

Raw 12-lead
(temporal)

VCG X,Y,Z
(spatial)

FFT magnitude
(spectral)

ResNet1d18 (shared backbone)
13 sigmoid outputs

confidence $s(x)$ from per-class decision margins

Learn-then-Test calibration on the target source
 $\max_{\tau} C(\tau)$ s.t. risk certified $\leq \alpha$

accept / abstain